



✓ Module 05 Lab - Data Preparation

Objective: To learn and apply the most common data preparation techniques. Raw data is rarely ready for a machine learning model. This process, also called preprocessing, is one of the most critical steps in the entire ML workflow.

In this lab, you will write more of the code. Read the explanations and then complete the tasks in the code cells.

✓ Part 1: Setup and Initial Look

We will continue using the Titanic dataset because it has the exact problems we need to solve: missing values and non-numeric data.

```
import pandas as pd
import numpy as np

# Load the dataset
df = pd.read_csv('https://raw.githubusercontent.com/datasciencedojo/datasets/master/titanic.csv')

# Let's look at the missing values
print("--- Missing Values Before Cleaning ---")
print(df.isnull().sum())
```

```
--- Missing Values Before Cleaning ---
PassengerId      0
Survived          0
Pclass           0
Name             0
Sex              0
Age             177
SibSp            0
Parch            0
Ticket           0
Fare             0
```

```
Cabin      687
Embarked    2
dtype: int64
```

✓ Part 2: Handling Missing Values (Imputation)

Concept: Most machine learning models cannot handle missing values (`NaN`). We must deal with them. Dropping the rows is an option, but you lose data. A better way is **imputation**, which means filling in the missing values with a calculated guess.

Common imputation strategies:

- **Mean:** Fill with the average value. Good for normally distributed data.
- **Median:** Fill with the middle value. Better for skewed data or data with outliers (like `Fare`).
- **Mode:** Fill with the most frequent value. Used for categorical data.

✓ Task 1: Impute the 'Age' Column

The 'Age' column is missing many values. Since age can be skewed (e.g., by a few very old passengers), using the **median** is a robust choice.

Your Task: Calculate the median of the 'Age' column and use the `.fillna()` method to replace the missing values.

```
# 1. Calculate the median of the 'Age' column
median_age = df['Age'].median()

# 2. Fill the missing values in 'Age' with the median
df['Age'].fillna(median_age, inplace=True)

# 3. Verify that there are no more missing values in 'Age'
print("Missing values in 'Age' after imputation:")
print(df['Age'].isnull().sum())
```

```
Missing values in 'Age' after imputation:
```

```
0
```

```
/tmp/ipykernel_559/2888799683.py:5: FutureWarning: A value is trying to be set on a copy of a DataFrame.  
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate
```

```
For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, i
```

```
df['Age'].fillna(median_age, inplace=True)
```

✓ Part 3: Encoding Categorical Features

Concept: Machine learning models are mathematical, so they need numbers, not text. We need to convert categorical columns (like 'Sex' and 'Embarked') into a numerical format. The most common method is **One-Hot Encoding**.

One-Hot Encoding takes a column with **N** categories and turns it into **N** new columns, each with a **1** or **0**. For example, the 'Sex' column (**male**, **female**) becomes two new columns: **Sex_male** and **Sex_female**.

Pandas has a convenient function called `pd.get_dummies()` that does this for us.

✓ Task 2: One-Hot Encode Categorical Columns

Your Task: Use `pd.get_dummies()` to encode the 'Sex' and 'Embarked' columns. Make sure to drop the original columns after encoding.

```
encoded_df = pd.get_dummies(df, columns=['Sex', 'Embarked'], drop_first=True)

print(encoded_df.head())
```

```
   PassengerId  Survived  Pclass  \
0             1         0       3
1             2         1       1
2             3         1       3
-
```

```

3          4          1          1
4          5          0          3

                                Name    Age  SibSp  Parch  \
0                                Braund, Mr. Owen Harris  22.0    1    0
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  38.0    1    0
2                                Heikkinen, Miss. Laina  26.0    0    0
3      Futrelle, Mrs. Jacques Heath (Lily May Peel)  35.0    1    0
4                                Allen, Mr. William Henry  35.0    0    0

                                Ticket    Fare  Cabin  Sex_male  Embarked_Q  Embarked_S
0      A/5 21171    7.2500  NaN    True    False    True
1      PC 17599   71.2833  C85   False    False    False
2  STON/O2. 3101282    7.9250  NaN   False    False    True
3      113803   53.1000  C123   False    False    True
4      373450    8.0500  NaN    True    False    True

```

✓ Part 4: Feature Scaling

Concept: Many models are sensitive to the scale of the features. For example, `Age` (from 0-80) and `Fare` (from 0-512) are on very different scales. This can cause the model to incorrectly believe that `Fare` is a more important feature simply because its values are larger.

Feature Scaling solves this by putting all features on a similar scale. A common method is **Standardization** (`StandardScaler` in scikit-learn), which rescales the data to have a mean of 0 and a standard deviation of 1.

Important: You only scale your numerical features, not your target variable or your newly encoded categorical columns.

✓ Task 3: Scale the 'Age' and 'Fare' Columns

Your Task: Use `StandardScaler` from `sklearn.preprocessing` to scale the 'Age' and 'Fare' columns.

```
from sklearn.preprocessing import StandardScaler
```

```
# --- ENTER YOUR CODE HERE ---
scaler = StandardScaler()

columns_to_scale = ['Age', 'Fare']

encoded_df[columns_to_scale] = scaler.fit_transform(encoded_df[columns_to_scale])

print(encoded_df.head())
```

```
   PassengerId  Survived  Pclass \
0             1         0       3
1             2         1       1
2             3         1       3
3             4         1       1
4             5         0       3
```

```
   Name                               Age  SibSp  Parch \
0  Braund, Mr. Owen Harris   -0.565736      1      0
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  0.663861      1      0
2  Heikkinen, Miss. Laina  -0.258337      0      0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)  0.433312      1      0
4  Allen, Mr. William Henry  0.433312      0      0
```

```
   Ticket    Fare  Cabin  Sex_male  Embarked_Q  Embarked_S
0  A/5 21171 -0.502445   NaN      True      False      True
1  PC 17599  0.786845   C85     False      False     False
2  STON/O2. 3101282 -0.488854   NaN     False      False     True
3  113803  0.420730  C123     False      False     True
4  373450 -0.486337   NaN      True      False     True
```



Knowledge Check

Instructions: Answer the following questions in this markdown cell.

1. Why is it often better to impute missing values with the median instead of the mean?
2. Explain in your own words what One-Hot Encoding does and why it is necessary.

3. Would you need to apply Feature Scaling to a Decision Tree model? Why or why not? (Hint: Think about how a Decision Tree makes its splits).

The median is often better because it is less affected by extremely high or low values, also known as outliers. Using the median can give a more accurate representation of the typical values in the dataset.

Explain in your own words what One-Hot Encoding does and why it is necessary.

With One-Hot Encoding, the text categories present in a column are taken and rewritten as separate columns for each category with True/False or 1/0 each options. Because of the use of `drop_first=True` command, one of the baseline categories is dropped in this case so that the information is not repeated multiple times. For example, the column can contain three values for Embarked; those three columns can be superseded by only two new columns, namely Embarked_Q and Embarked_S. A person with False in both columns must have boarded in the category that was eliminated. The process is necessary due to the fact that machine learning models rely on mathematical operations and cannot deal with text, thus we use this type of encoding to make sure that the model does not make false assumptions about unranked categories.

Would you need to apply Feature Scaling to a Decision Tree model? Why or why not? (Hint: Think about how a Decision Tree makes its splits).

No, it is unnecessary to use feature scaling with a model based on decision trees. When building a decision tree, the only thing that the algorithm does is search for threshold values from a particular feature (for example, whether Age > 30 or Fare < 50). Since the model does not measure the distance between data points or uses any kind of gradient descent to optimize the algorithm, the scale of numbers does not matter. Feature scaling is important for the algorithms that react to distances and scales only, which includes KNN, SVM, K-Means, and Neural Networks.

